



Published models on our 177 held-out clips



Which alphabet each model answered in



Every model decodes under its own generation config; no language or task token is imposed on anyone else's model. Our references are romanized, so a Bengali-script answer scores near 100 per cent however well the model heard the speech. The orange bars transliterate that output to Roman and keep the better of two schemes per clip, so they are an upper bound on what each model could score if orthography were free.